

prabhudayala@gmail.com_5

April 24, 2019

1 DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result

How to scale current manual processes and resources to screen 500,000 projects so that they can
How to increase the consistency of project vetting across different volunteers to improve t
How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

1.1 About the DonorsChoose Data Set

The `train.csv` data set provided by DonorsChoose contains the following features:

Feature	Description
<code>project_id</code>	A unique identifier for the proposed project. Example: p036502

`project_title` | Title of the project. **Examples:**

Art Will Make You Happy!

First Grade Fun

`project_grade_category` | Grade level of students for which the project is targeted. One of the following enumerated values:

Grades PreK-2

Grades 3-5

Grades 6-8

Grades 9-12

`project_subject_categories` | One or more (comma-separated) subject categories for the project from the following enumerated list of values:

Applied Learning
Care & Hunger
Health & Sports
History & Civics
Literacy & Language
Math & Science
Music & The Arts
Special Needs
Warmth

Examples:

Music & The Arts
Literacy & Language, Math & Science

school_state | State where school is located ([Two-letter U.S. postal code](#)). **Example:** WY
project_subject_subcategories | One or more (comma-separated) subject subcategories for the project. **Examples:**

Literacy
Literature & Writing, Social Sciences

project_resource_summary | An explanation of the resources needed for the project. **Example:**

My students need hands on literacy materials to manage sensory needs!

project_essay_1 | First application essay

project_essay_2 | *Second application essay* **project_essay_3** | Third application essay

project_essay_4 | *Fourth application essay* **project_submitted_datetime** | Datetime when project application was submitted. **Example:** 2016-04-28 12:43:56.245

teacher_id | A unique identifier for the teacher of the proposed project. **Example:** bdf8baa8fedef6bfeec7ae4ff1c15c56

teacher_prefix | Teacher's title. One of the following enumerated values:

nan
Dr.
Mr.
Mrs.
Ms.
Teacher.

teacher_number_of_previously_posted_projects | Number of project applications previously submitted by the same teacher. **Example:** 2

* See the section Notes on the Essay Data for more details about these features.

Additionally, the `resources.csv` data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
id	A <code>project_id</code> value from the <code>train.csv</code> file. Example: p036502

Feature	Description
description	Description of the resource. Example: Tenor Saxophone Reeds, Box of 25
quantity	Quantity of the resource required. Example: 3
price	Price of the resource required. Example: 9.95

Note: Many projects require multiple resources. The `id` value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
project_is_approved	Adjudication flag indicating whether DonorsChoose approved the project. A value of 0 indicates the project was not approved, and a value of 1 indicates the project was approved.

1.1.1 Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

project_essay_1: "Introduce us to your classroom"

project_essay_2: "Tell us more about your students"

project_essay_3: "Describe how your students will use the materials you're requesting"

project_essay_3: "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

project_essay_1: "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."

project_essay_2: "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be NaN.

```
In [1]: %matplotlib inline
import warnings
```

```
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from plotly import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter
```

1.2 1.1 Reading Data

```
In [2]: project_data = pd.read_csv('train_data.csv')
        resource_data = pd.read_csv('resources.csv')

In [3]: print("Number of data points in train data", project_data.shape)
        print('-'*50)
        print("The attributes of data :", project_data.columns.values)
```

Number of data points in train data (109248, 17)

```
The attributes of data : ['Unnamed: 0' 'id' 'teacher_id' 'teacher_prefix' 'school_state'
'project_submitted_datetime' 'project_grade_category'
'project_subject_categories' 'project_subject_subcategories'
'project_title' 'project_essay_1' 'project_essay_2' 'project_essay_3'
'project_essay_4' 'project_resource_summary'
'teacher_number_of_previously_posted_projects' 'project_is_approved']
```

```
In [4]: print("Number of data points in train data", resource_data.shape)
        print(resource_data.columns.values)
        resource_data.head(2)
```

```
Number of data points in train data (1541272, 4)
['id' 'description' 'quantity' 'price']
```

```
Out [4]:
```

	id	description	quantity	\
0	p233245	LC652 - Lakeshore Double-Space Mobile Drying Rack	1	
1	p069063	Bouncy Bands for Desks (Blue support pipes)	3	

	price
0	149.00
1	14.95

```
In [5]: # join two dataframes in python:
        price_data = resource_data.groupby('id').agg({'price': 'sum', 'quantity': 'sum'}).reset_index()
        price_data.head(2)
        project_data = pd.merge(project_data, price_data, on='id', how='left')
```

```
In [6]: project_data.columns
```

```
Out [6]: Index(['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',
'project_submitted_datetime', 'project_grade_category',
'project_subject_categories', 'project_subject_subcategories',
'project_title', 'project_essay_1', 'project_essay_2',
'project_essay_3', 'project_essay_4', 'project_resource_summary',
'teacher_number_of_previously_posted_projects', 'project_is_approved',
'price', 'quantity'],
dtype='object')
```

1.3 1.2 preprocessing of project_subject_categories

```
In [7]: categories = list(project_data['project_subject_categories'].values)
        # remove special characters from list of strings python: https://stackoverflow.com/a/4758804

        # https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
        # https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
        # https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python
        cat_list = []
```

```

for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space
            j=j.replace('The','') # if we have the words "The" we are going to replace
            j = j.replace(' ','') # we are replacing all the ' '(space) with ''(empty) ex: "Math & Science"
            temp+=j.strip()+" " # " abc ".strip() will return "abc", remove the trailing space
            temp = temp.replace('&','_') # we are replacing the & value into _
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))

```

1.4 1.3 preprocessing of project_subject_subcategories

```

In [8]: sub_categories = list(project_data['project_subject_subcategories'].values)
        # remove special characters from list of strings python: https://stackoverflow.com/a/40000000

        # https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
        # https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
        # https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

sub_cat_list = []
for i in sub_categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space
            j=j.replace('The','') # if we have the words "The" we are going to replace
            j = j.replace(' ','') # we are replacing all the ' '(space) with ''(empty) ex: "Math & Science"
            temp +=j.strip()+" " # " abc ".strip() will return "abc", remove the trailing space
            temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/408403
my_counter = Counter()

```

```

for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

1.5 1.3 Text preprocessing

In [9]: # merge two column text dataframe:

```

project_data["essay"] = project_data["project_essay_1"].map(str) + \
    project_data["project_essay_2"].map(str) + \
    project_data["project_essay_3"].map(str) + \
    project_data["project_essay_4"].map(str)

```

In [10]: project_data.head(2)

```

Out[10]:   Unnamed: 0      id      teacher_id teacher_prefix \
0      160221  p253737  c90749f5d961ff158d4b4d1e7dc665fc  Mrs.
1      140945  p258326  897464ce9ddc600bcd1151f324dd63a    Mr.

  school_state project_submitted_datetime project_grade_category \
0           IN      2016-12-05 13:43:57      Grades PreK-2
1           FL      2016-10-25 09:22:10      Grades 6-8

  project_title \
0  Educational Support for English Learners at Home
1           Wanted: Projector for Hungry Learners

  project_essay_1 \
0  My students are English learners that are work...
1  Our students arrive to our school eager to lea...

  project_essay_2 project_essay_3 \
0  \"The limits of your language are the limits o...      NaN
1  The projector we need for our school is very c...      NaN

  project_essay_4      project_resource_summary \
0           NaN  My students need opportunities to practice beg...
1           NaN  My students need a projector to help with view...

  teacher_number_of_previously_posted_projects  project_is_approved  price \
0                                           0                      0  154.6
1                                           7                      1  299.0

  quantity      clean_categories      clean_subcategories \
0       23      Literacy_Language      ESL Literacy
1       1  History_Civics Health_Sports  Civics_Government TeamSports

```

```

                                essay
0  My students are English learners that are work...
1  Our students arrive to our school eager to lea...

```

In [11]: *#### 1.4.2.3 Using Pretrained Models: TFIDF weighted W2V*

```

In [12]: # printing some random reviews
print(project_data['essay'].values[0])
print("="*50)
print(project_data['essay'].values[150])
print("="*50)
print(project_data['essay'].values[1000])
print("="*50)
print(project_data['essay'].values[20000])
print("="*50)
print(project_data['essay'].values[99999])
print("="*50)

```

```

My students are English learners that are working on English as their second or third language
=====
The 51 fifth grade students that will cycle through my classroom this year all love learning, a
=====
How do you remember your days of school? Was it in a sterile environment with plain walls, row
=====
My kindergarten students have varied disabilities ranging from speech and language delays, cog
=====
The mediocre teacher tells. The good teacher explains. The superior teacher demonstrates. The g
=====

```

In [13]: *# <https://stackoverflow.com/a/47091490/4084039>*

```

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

```



```
In [14]: sent = decontracted(project_data['essay'].values[20000])
        print(sent)
        print("="*50)
```

My kindergarten students have varied disabilities ranging from speech and language delays, cog
=====

```
In [15]: # \r \n \t remove from string python: http://texthandler.com/info/remove-line-breaks-
        sent = sent.replace('\r', ' ')
        sent = sent.replace('\n', ' ')
        sent = sent.replace('\t', ' ')
        print(sent)
```

My kindergarten students have varied disabilities ranging from speech and language delays, cog

```
In [16]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
        sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
        print(sent)
```

My kindergarten students have varied disabilities ranging from speech and language delays cog

```
In [17]: # https://gist.github.com/sebleier/554280
        # we are removing the words from the stop words list: 'no', 'nor', 'not'
        stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you'
            "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him'
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
            'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'h
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as
            'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'o
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'an
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", '
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't
            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi
            "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
            'won', "won't", 'wouldn', "wouldn't"]
```

```
In [18]: # Combining all the above stundents
        from tqdm import tqdm
        preprocessed_essays = []
        # tqdm is for printing the status bar
        for sentence in tqdm(project_data['essay'].values):
            sent = decontracted(sentence)
            sent = sent.replace('\r', ' ')
```

```

sent = sent.replace('\\\"', ' ')
sent = sent.replace('\\n', ' ')
sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
# https://gist.github.com/sebleier/554280
sent = ' '.join(e for e in sent.split() if e not in stopwords)
preprocessed_essays.append(sent.lower().strip())

```

100%|| 109248/109248 [00:44<00:00, 2459.88it/s]

```

In [19]: # after preprocessing
preprocessed_essays[20000]

```

Out[19]: 'my kindergarten students varied disabilities ranging speech language delays cognitive'

```

In [20]: project_data["essay"]=preprocessed_essays
project_data.essay.values[20000]

```

Out[20]: 'my kindergarten students varied disabilities ranging speech language delays cognitive'

1.4 Preprocessing of project_title

```

In [21]: from tqdm import tqdm
preprocessed_project_title = []
# tqdm is for printing the status bar
for sentence in tqdm(project_data['project_title'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\\"', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
    preprocessed_project_title.append(sent.lower().strip())

```

100%|| 109248/109248 [00:01<00:00, 57298.53it/s]

```

In [22]: print(project_data['project_title'].values[20000])
project_data['project_title']=preprocessed_project_title
print(project_data['project_title'].values[20000])

```

We Need To Move It While We Input It!
need move input

```

In [23]: from nltk.sentiment import SentimentIntensityAnalyzer as SID
#nltk.download('vader_lexicon')
new_df_as_dictionary=[]
sid=SID()
for i in tqdm(project_data.essay.values):
    new_df_as_dictionary.append(sid.polarity_scores(i))

```

100%|| 109248/109248 [02:23<00:00, 758.93it/s]

```
In [24]: print(project_data.columns)
         print(project_data.shape)
         sentiment_score=pd.DataFrame(new_df_as_dictionary)
         print(sentiment_score.columns)
         print(sentiment_score.shape)
```

```
Index(['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',
       'project_submitted_datetime', 'project_grade_category', 'project_title',
       'project_essay_1', 'project_essay_2', 'project_essay_3',
       'project_essay_4', 'project_resource_summary',
       'teacher_number_of_previously_posted_projects', 'project_is_approved',
       'price', 'quantity', 'clean_categories', 'clean_subcategories',
       'essay'],
      dtype='object')
(109248, 20)
Index(['compound', 'neg', 'neu', 'pos'], dtype='object')
(109248, 4)
```

```
In [25]: sentiment_score=pd.DataFrame(new_df_as_dictionary)
         project_data=pd.concat((project_data,sentiment_score),axis=1,ignore_index=True)
         print(project_data.shape)

(109248, 24)
```

```
In [26]: project_data.columns=['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',
                               'project_submitted_datetime', 'project_grade_category', 'project_title',
                               'project_essay_1', 'project_essay_2', 'project_essay_3',
                               'project_essay_4', 'project_resource_summary',
                               'teacher_number_of_previously_posted_projects', 'project_is_approved',
                               'price', 'quantity', 'clean_categories', 'clean_subcategories',
                               'essay', 'compound', 'neg', 'neu', 'pos']
```

```
In [27]: print(project_data.head(5))
```

	Unnamed: 0	id	teacher_id	teacher_prefix	\
0	160221	p253737	c90749f5d961ff158d4b4d1e7dc665fc	Mrs.	
1	140945	p258326	897464ce9ddc600bcd1151f324dd63a	Mr.	
2	21895	p182444	3465aaf82da834c0582ebd0ef8040ca0	Ms.	
3	45	p246581	f3cb9bffbba169bef1a77b243e620b60	Mrs.	
4	172407	p104768	be1f7507a41f8479dc06f047086a39ec	Mrs.	

	school_state	project_submitted_datetime	project_grade_category	\
0	IN	2016-12-05 13:43:57	Grades PreK-2	
1	FL	2016-10-25 09:22:10	Grades 6-8	

2	AZ	2016-08-31 12:03:56	Grades 6-8
3	KY	2016-10-06 21:16:17	Grades PreK-2
4	TX	2016-07-11 01:10:09	Grades PreK-2

```

                                project_title \
0      educational support english learners home
1              wanted projector hungry learners
2 soccer equipment awesome middle school students
3              techie kindergarteners
4              interactive math tools

```

```

                                project_essay_1 \
0 My students are English learners that are work...
1 Our students arrive to our school eager to lea...
2 \r\n\"True champions aren't always the ones th...
3 I work at a unique school filled with both ESL...
4 Our second grade classroom next year will be m...

```

```

                                project_essay_2 ... \
0 \"The limits of your language are the limits o... ...
1 The projector we need for our school is very c... ...
2 The students on the campus come to school know... ...
3 My students live in high poverty conditions wi... ...
4 For many students, math is a subject that does... ...

```

	project_is_approved	price	quantity	clean_categories
0	0	154.60	23	Literacy_Language
1	1	299.00	1	History_Civics Health_Sports
2	0	516.85	22	Health_Sports
3	1	232.90	4	Literacy_Language Math_Science
4	1	67.98	4	Math_Science

```

                                clean_subcategories \
0      ESL Literacy
1 Civics_Government TeamSports
2 Health_Wellness TeamSports
3 Literacy Mathematics
4 Mathematics

```

	essay	compound	neg	neu
0 my students english learners working english s...	0.9694	0.012	0.844	
1 our students arrive school eager learn they po...	0.9856	0.048	0.669	
2 true champions not always ones win guts by mia...	0.9816	0.122	0.659	
3 i work unique school filled esl english second...	0.9656	0.106	0.649	
4 our second grade classroom next year made arou...	0.8524	0.066	0.791	

```

pos
0 0.144

```

```

1  0.283
2  0.219
3  0.246
4  0.143

```

[5 rows x 24 columns]

2 Assignment 5: Logistic Regression

```

<li><strong>[Task-1] Logistic Regression(either SGDClassifier with log loss, or LogisticRegression)
    <ul>
        <li><font color='red'>Set 1</font>: categorical, numerical features + project_title(BOU
        <li><font color='red'>Set 2</font>: categorical, numerical features + project_title(TF
        <li><font color='red'>Set 3</font>: categorical, numerical features + project_title(AV
        <li><font color='red'>Set 4</font>: categorical, numerical features + project_title(TF
    </li>
<br>
<li><strong>Hyper paramter tuning (find best hyper parameters corresponding the algorithm that
    <ul>
<li>Find the best hyper parameter which will give the maximum <a href='https://www.appliedaicom
<li>Find the best hyper paramter using k-fold cross validation or simple cross validation data
<li>Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this ta
    </ul>
</li>
<br>
<li><strong>Representation of results</strong>
    <ul>
<li>You need to plot the performance of model both on train data and cross validation data for
<img src='train_cv_auc.JPG' width=300px></li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and f
<img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicomse.
<img src='confusion_matrix.png' width=300px></li>
    </ul>
</li>
<br>
<li><strong>[Task-2] Apply Logistic Regression on the below feature set <font color='red'> Set
<li> Consider these set of features <font color='red'> Set 5 :</font>
    <ul>
        <li><strong>school_state</strong> : categorical data</li>
        <li><strong>clean_categories</strong> : categorical data</li>
        <li><strong>clean_subcategories</strong> : categorical data</li>
        <li><strong>project_grade_category</strong> :categorical data</li>
        <li><strong>teacher_prefix</strong> : categorical data</li>
        <li><strong>quantity</strong> : numerical data</li>
        <li><strong>teacher_number_of_previously_posted_projects</strong> : numerical data

```

```

        <li><strong>price</strong> : numerical data</li>
        <li><strong>sentiment score's of each of the essay</strong> : numerical data</li>
        <li><strong>number of words in the title</strong> : numerical data</li>
        <li><strong>number of words in the combine essays</strong> : numerical data</li>
    </ul>
    And apply the Logistic regression on these features by finding the best hyper paramter as :
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
        <img src='summary.JPG' width=400px>
</li>
    </ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

2. Logistic Regression

2.1 Splitting data into Train and cross validation(or test): Stratified Sampling

```

In [28]: sampling=False
undersampling=False
if (not sampling):
    print("Total data ",project_data.shape)

else:
    if(sampling and undersampling):
        print("Total data ",project_data.shape)
        project_data_negative=project_data[project_data.project_is_approved==0]
        project_data_positive=project_data[project_data.project_is_approved==1]
        project_data_positive=project_data_positive.sample(n=project_data_negative.shape[0])
        print("Positive points: ",project_data_positive.shape[0])
        print("Negaitive points: ",project_data_negative.shape[0])
        project_data=pd.concat([project_data_positive,project_data_negative])
    else:
        print("Total data ",project_data.shape)
        project_data_negative=project_data[project_data.project_is_approved==0]
        project_data_positive=project_data[project_data.project_is_approved==1]
        project_data_negative=project_data_negative.sample(n=project_data_positive.shape[0])
        print("Positive points: ",project_data_positive.shape[0])
        print("Negaitive points: ",project_data_negative.shape[0])

```

```

project_data=pd.concat([project_data_positive,project_data_negative])

#data_point_size=50000
#project_data=project_data.sample(n=data_point_size,random_state=42,replace=True)
print("positive and negative counts")
print(project_data.project_is_approved.value_counts())
project_data_Y=project_data.project_is_approved
project_data_X=project_data.drop(columns=['project_is_approved'])
print("After sampling: ",project_data_X.shape)

Total data (109248, 24)
positive and negative counts
1    92706
0    16542
Name: project_is_approved, dtype: int64
After sampling: (109248, 23)

In [29]: from sklearn.model_selection import train_test_split
project_data_X_train,project_data_X_test,project_data_Y_train,project_data_Y_test=train_test_split(project_data_X,project_data_Y,train_size=0.8,random_state=42)

In [30]: print(project_data_X_train.shape)
print(project_data_X_test.shape)
print(project_data_Y_train.shape)
print(project_data_Y_test.shape)

(87398, 23)
(21850, 23)
(87398,)
(21850,)

```

2.2 Make Data Model Ready: encoding numerical, categorical features

2.2.1 Categorical features

```

In [31]: from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(vocabulary=list(sorted_cat_dict.keys()), lowercase=False)
vectorizer.fit(project_data_X_train['clean_categories'].values)
print(vectorizer.get_feature_names())

#for train data
categories_one_hot_train = vectorizer.transform(project_data_X_train['clean_categories'].values)
print("Shape of matrix after one hot encodig ",categories_one_hot_train.shape)

#for test
categories_one_hot_test = vectorizer.transform(project_data_X_test['clean_categories'].values)
print("Shape of matrix after one hot encodig ",categories_one_hot_test.shape)

```

```
['Warmth', 'Care_Hunger', 'History_Civics', 'Music_Arts', 'AppliedLearning', 'SpecialNeeds', '']
Shape of matrix after one hot encodig (87398, 9)
Shape of matrix after one hot encodig (21850, 9)
```

```
In [32]: vectorizer = CountVectorizer(vocabulary=list(sorted_sub_cat_dict.keys()), lowercase=False)
vectorizer.fit(project_data_X_train['clean_subcategories'].values)
print(vectorizer.get_feature_names())
```

```
#for train data
```

```
sub_categories_one_hot_train = vectorizer.transform(project_data_X_train['clean_subcategories'].values)
print("Shape of matrix after one hot encodig ",sub_categories_one_hot_train.shape)
```

```
#for test
```

```
sub_categories_one_hot_test = vectorizer.transform(project_data_X_test['clean_subcategories'].values)
print("Shape of matrix after one hot encodig ",sub_categories_one_hot_test.shape)
```

```
['Economics', 'CommunityService', 'FinancialLiteracy', 'ParentInvolvement', 'Extracurricular', '']
Shape of matrix after one hot encodig (87398, 30)
Shape of matrix after one hot encodig (21850, 30)
```

```
In [33]: project_data_X_train.teacher_prefix = project_data_X_train.teacher_prefix.replace(np.nan, 'Mrs.')
print(project_data_X_train.teacher_prefix.value_counts())
project_data_X_test.teacher_prefix = project_data_X_test.teacher_prefix.replace(np.nan, 'Mrs.')
print(project_data_X_test.teacher_prefix.value_counts())
```

```
Mrs.      45800
```

```
Ms.       31168
```

```
Mr.       8519
```

```
Teacher   1898
```

```
Dr.       11
```

```
2
```

```
Name: teacher_prefix, dtype: int64
```

```
Mrs.      11469
```

```
Ms.       7787
```

```
Mr.       2129
```

```
Teacher   462
```

```
Dr.       2
```

```
1
```

```
Name: teacher_prefix, dtype: int64
```

```
In [34]: # we use count vectorizer to convert the values into one hot encoded features
vectorizer = CountVectorizer(vocabulary=['Mrs.', 'Ms.', 'Mr.', 'Teacher', 'Dr.'], lowercase=False)
vectorizer.fit(project_data_X_train['teacher_prefix'].values)
print(vectorizer.get_feature_names())
```

```
teacher_prefix_one_hot_train = vectorizer.transform(project_data_X_train['teacher_prefix'].values)
```



```

print("Shape of matrix after one hot encodig ",teacher_prefix_one_hot_train.shape)

teacher_prefix_one_hot_test = vectorizer.transform(project_data_X_test['teacher_prefix'])
print("Shape of matrix after one hot encodig ",teacher_prefix_one_hot_test.shape)

['Mrs.', 'Ms.', 'Mr.', 'Teacher', 'Dr.']
Shape of matrix after one hot encodig  (87398, 5)
Shape of matrix after one hot encodig  (21850, 5)

In [35]: # we use count vectorizer to convert the values into one hot encoded features
vectorizer = CountVectorizer(vocabulary=list(project_data_X_train['project_grade_category'].values))
vectorizer.fit(project_data_X_train['project_grade_category'].values)
print(vectorizer.get_feature_names())

project_grade_category_one_hot_train = vectorizer.transform(project_data_X_train['project_grade_category'])
print("Shape of matrix after one hot encodig ",project_grade_category_one_hot_train.shape)

project_grade_category_one_hot_test = vectorizer.transform(project_data_X_test['project_grade_category'])
print("Shape of matrix after one hot encodig ",project_grade_category_one_hot_test.shape)

['Grades PreK-2', 'Grades 3-5', 'Grades 6-8', 'Grades 9-12']
Shape of matrix after one hot encodig  (87398, 4)
Shape of matrix after one hot encodig  (21850, 4)

In [36]: # we use count vectorizer to convert the values into one hot encoded features
vectorizer = CountVectorizer(vocabulary=list(project_data_X_train['school_state'].unique().values))
vectorizer.fit(project_data_X_train['school_state'].values)
print(vectorizer.get_feature_names())

school_state_one_hot_train = vectorizer.transform(project_data_X_train['school_state'])
print("Shape of matrix after one hot encodig ",school_state_one_hot_train.shape)

school_state_one_hot_test = vectorizer.transform(project_data_X_test['school_state'])
print("Shape of matrix after one hot encodig ",school_state_one_hot_test.shape)

['NY', 'MD', 'OK', 'MA', 'CA', 'AR', 'FL', 'PA', 'SC', 'NC', 'AZ', 'MI', 'AL', 'WI', 'NV', 'UT']
Shape of matrix after one hot encodig  (87398, 51)
Shape of matrix after one hot encodig  (21850, 51)

```

2.2.2 Numerical features

```

In [37]: # check this one: https://www.youtube.com/watch?v=0HQqDcln3Z4&t=530s
# standardization sklearn: https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler
from sklearn.preprocessing import StandardScaler

```

```

# price_standardized = standardScaler.fit(project_data['price'].values)
# this will rise the error
# ValueError: Expected 2D array, got 1D array instead: array=[725.05 213.03 329.    .
# Reshape your data either using array.reshape(-1, 1)

price_scalar = StandardScaler()
price_scalar.fit(project_data_X_train['price'].values.reshape(-1,1)) # finding the me
print(f"Mean : {price_scalar.mean_[0]}, Standard deviation : {np.sqrt(price_scalar.va

# Now standardize the data with above maen and variance.
price_standardized_train = price_scalar.transform(project_data_X_train['price'].values
# Now standardize the data with above maen and variance.
price_standardized_test = price_scalar.transform(project_data_X_test['price'].values.

```

Mean : 298.64356770177807, Standard deviation : 368.42853396795914

```

In [38]: # check this one: https://www.youtube.com/watch?v=0HQqOcln3Z4&t=530s
# standardization sklearn: https://scikit-learn.org/stable/modules/generated/sklearn.
from sklearn.preprocessing import StandardScaler,normalize

# price_standardized = standardScaler.fit(project_data['price'].values)
# this will rise the error
# ValueError: Expected 2D array, got 1D array instead: array=[725.05 213.03 329.    .
# Reshape your data either using array.reshape(-1, 1)

price_scalar = StandardScaler()
price_scalar.fit(project_data_X_train['teacher_number_of_previously_posted_projects']
print(f"Mean : {price_scalar.mean_[0]}, Standard deviation : {np.sqrt(price_scalar.va

# Now standardize the data with above maen and variance.
teacher_number_of_previously_posted_projects_standardized_train = price_scalar.transf
# Now standardize the data with above maen and variance.
teacher_number_of_previously_posted_projects_standardized_test = price_scalar.transfor

```

Mean : 11.102897091466623, Standard deviation : 27.572082372998246

2.3 Make Data Model Ready: encoding eassay, and project_title

```

In [39]: vectorizer = CountVectorizer(min_df=10,ngram_range=(1,2),max_features=5000)
vectorizer.fit(project_data_X_train.essay.values)

text_bow_train=vectorizer.fit_transform(project_data_X_train.essay.values)
print(text_bow_train.shape)

text_bow_test=vectorizer.transform(project_data_X_test.essay.values)
print(text_bow_test.shape)

```

```
(87398, 5000)
(21850, 5000)
```

```
In [40]: # Similarly you can vectorize for title also
vectorizer = CountVectorizer(min_df=10)
vectorizer.fit(project_data_X_train.project_title.values)

title_text_bow_train=vectorizer.fit_transform(project_data_X_train.project_title.values)
print(title_text_bow_train.shape)

title_text_bow_test=vectorizer.transform(project_data_X_test.project_title.values)
print(title_text_bow_test.shape)
```

```
(87398, 2803)
(21850, 2803)
```

```
In [41]: from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=10,ngram_range=(1,2),max_features=5000)
vectorizer.fit(project_data_X_train.essay.values)

text_tfidf_train=vectorizer.fit_transform(project_data_X_train.essay.values)
print(text_tfidf_train.shape)

text_tfidf_test=vectorizer.transform(project_data_X_test.essay.values)
print(text_tfidf_test.shape)
```

```
(87398, 5000)
(21850, 5000)
```

```
In [42]: # Similarly you can vectorize for title also
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=10)
vectorizer.fit(project_data_X_train.project_title.values)

title_text_tfidf_train=vectorizer.fit_transform(project_data_X_train.project_title.values)
print(title_text_tfidf_train.shape)

title_text_tfidf_test=vectorizer.transform(project_data_X_test.project_title.values)
print(title_text_tfidf_test.shape)
```

```
(87398, 2803)
(21850, 2803)
```

```
In [43]: # Reading glove vectors in python: https://stackoverflow.com/a/38230349/4084039
def loadGloveModel(gloveFile):
```

```

print ("Loading Glove Model")
f = open(gloveFile, 'r', encoding="utf8")
model = {}
for line in tqdm(f):
    splitLine = line.split()
    word = splitLine[0]
    embedding = np.array([float(val) for val in splitLine[1:]])
    model[word] = embedding
print ("Done.", len(model), " words loaded!")
return model
# borrowed from https://therenegadecoder.com/code/how-to-check-if-a-file-exists-in-py
import os
exists = os.path.isfile('./glove_vectors')
if(not exists):
    model = loadGloveModel('glove.42B.300d.txt')

'''# =====
Output:

Loading Glove Model
1917495it [06:32, 4879.69it/s]
Done. 1917495 words loaded!

# ====='''

words = []
for i in preproced_texts:
    words.extend(i.split(' '))

for i in preproced_titles:
    words.extend(i.split(' '))
print("all the words in the coupus", len(words))
words = set(words)
print("the unique words in the coupus", len(words))

inter_words = set(model.keys()).intersection(words)
print("The number of words that are present in both glove vectors and our coupus"
      len(inter_words), "(", np.round(len(inter_words)/len(words)*100, 3), "%)")

words_courpus = {}
words_glove = set(model.keys())
for i in words:
    if i in words_glove:
        words_courpus[i] = model[i]
print("word 2 vec length", len(words_courpus))

# stronging variables into pickle files python: http://www.jessicayung.com/how-to

```

```

import pickle
with open('glove_vectors', 'wb') as f:
    pickle.dump(words_courpus, f)
else:
    print("glove already exists. No need to compute")

```

glove already exists. No need to compute

```

In [44]: with open('glove_vectors', 'rb') as f:
        model = pickle.load(f)
        glove_words = set(model.keys())

```

```

In [45]: # average Word2Vec
        # compute average word2vec for each review.
avg_w2v_vectors_essay_train = []; # the avg-w2v for each sentence/review is stored in
for sentence in tqdm(project_data_X_train.essay.values): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_vectors_essay_train.append(vector)

print(len(avg_w2v_vectors_essay_train))
print(len(avg_w2v_vectors_essay_train[0]))

```

100%|| 87398/87398 [00:19<00:00, 4560.29it/s]

87398

300

```

In [46]: # average Word2Vec
        # compute average word2vec for each review.
avg_w2v_vectors_essay_test = []; # the avg-w2v for each sentence/review is stored in
for sentence in tqdm(project_data_X_test.essay.values): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:

```

```

        vector /= cnt_words
        avg_w2v_vectors_essay_test.append(vector)

print(len(avg_w2v_vectors_essay_test))
print(len(avg_w2v_vectors_essay_test[0]))

```

100%|| 21850/21850 [00:04<00:00, 4514.38it/s]

21850
300

```

In [47]: # average Word2Vec
# compute average word2vec for each title.
title_avg_w2v_vectors_train = []; # the avg-w2v for each sentence/review is stored in
for sentence in tqdm(project_data_X_train.project_title.values): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    title_avg_w2v_vectors_train.append(vector)

print(len(title_avg_w2v_vectors_train))
print(len(title_avg_w2v_vectors_train[0]))

```

100%|| 87398/87398 [00:00<00:00, 88824.17it/s]

87398
300

```

In [48]: # average Word2Vec
# compute average word2vec for each title.
title_avg_w2v_vectors_test = []; # the avg-w2v for each sentence/review is stored in
for sentence in tqdm(project_data_X_test.project_title.values): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words

```

```

        title_avg_w2v_vectors_test.append(vector)

    print(len(title_avg_w2v_vectors_test))
    print(len(title_avg_w2v_vectors_test[0]))

```

100%|| 21850/21850 [00:00<00:00, 88577.91it/s]

21850
300

```

In [49]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
tfidf_model = TfidfVectorizer()
tfidf_model.fit(project_data_X_train.essay.values)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
essay_tfidf_words = set(tfidf_model.get_feature_names())

```

```

In [50]: # average Word2Vec
# compute average word2vec for each review.
essay_tfidf_w2v_vectors_train = []; # the avg-w2v for each sentence/review is stored
for sentence in tqdm(project_data_X_train.essay.values): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in essay_tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value((s
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) #
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    essay_tfidf_w2v_vectors_train.append(vector)

    print(len(essay_tfidf_w2v_vectors_train))
    print(len(essay_tfidf_w2v_vectors_train[0]))

```

100%|| 87398/87398 [02:17<00:00, 637.26it/s]

87398
300

```

In [51]: # average Word2Vec
# compute average word2vec for each review.
essay_tfidf_w2v_vectors_test = []; # the avg-w2v for each sentence/review is stored i

```

```

for sentence in tqdm(project_data_X_test.essay.values): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in essay_tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value((s
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) #
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    essay_tfidf_w2v_vectors_test.append(vector)

print(len(essay_tfidf_w2v_vectors_test))
print(len(essay_tfidf_w2v_vectors_test[0]))

```

100%|| 21850/21850 [00:35<00:00, 617.76it/s]

21850

300

```

In [52]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
tfidf_model = TfidfVectorizer()
tfidf_model.fit(project_data_X_train.project_title.values)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
essay_tfidf_words = set(tfidf_model.get_feature_names())

```

```

In [53]: # average Word2Vec
# compute average word2vec for each review.
title_tfidf_w2v_vectors_train = []; # the avg-w2v for each sentence/review is stored
for sentence in tqdm(project_data_X_train.project_title.values): # for each review/se
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in essay_tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value((s
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) #
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    title_tfidf_w2v_vectors_train.append(vector)

```



```

print(len(title_tfidf_w2v_vectors_train))
print(len(title_tfidf_w2v_vectors_train[0]))
100%|| 87398/87398 [00:01<00:00, 44042.37it/s]

87398
300

In [54]: # average Word2Vec
# compute average word2vec for each review.
title_tfidf_w2v_vectors_test = []; # the avg-w2v for each sentence/review is stored in
for sentence in tqdm(project_data_X_test.project_title.values): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in essay_tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf value(sentence.count(word)/len(sentence.split()))
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # calculating tfidf weighted w2v
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    title_tfidf_w2v_vectors_test.append(vector)

print(len(title_tfidf_w2v_vectors_test))
print(len(title_tfidf_w2v_vectors_test[0]))
100%|| 21850/21850 [00:00<00:00, 46247.65it/s]

21850
300

```

2.4 Applying Logistic Regression on different kind of featurization as mentioned in the instructions

Apply Logistic Regression on different kind of featurization as mentioned in the instructions
For Every model that you work on make sure you do the step 2 and step 3 of instructions

2.0.1 2.4.1 Applying LR on BOW, SET 1

```

In [55]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack
# with the same hstack function we are concatenating a sparse matrix and a dense matrix
BOW = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_state_one_hot_train))
print(BOW.shape)
BOW_test = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_state_one_hot_test))
print(BOW_test.shape)

```

```
(87398, 7900)
(21850, 7900)
```

```
In [56]: from sklearn.linear_model import LogisticRegression
         from sklearn.model_selection import GridSearchCV
         model=LogisticRegression(class_weight='balanced')
         a=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
         print(a)
         parameters = {'C': a }
         clf = GridSearchCV(model, parameters, cv=4,scoring='roc_auc',n_jobs=4,verbose=10)
         clf.fit(BOW,project_data_Y_train)
```

```
[0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]
```

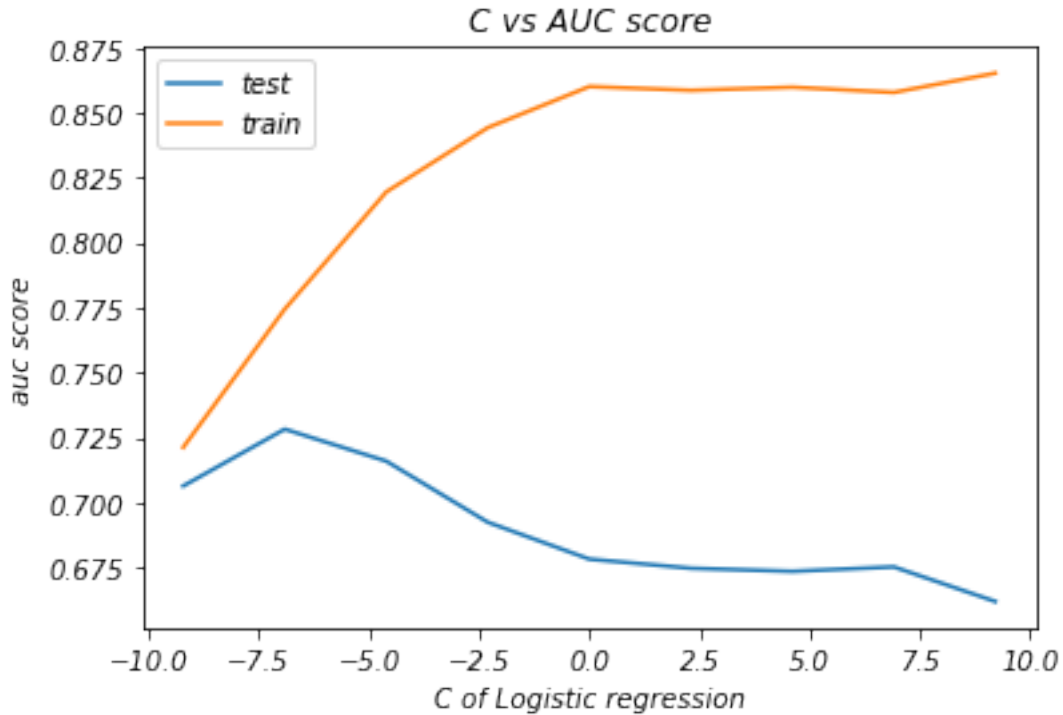
Fitting 4 folds for each of 9 candidates, totalling 36 fits

```
[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done   5 tasks      | elapsed:   18.7s
[Parallel(n_jobs=4)]: Done  10 tasks      | elapsed:   40.3s
[Parallel(n_jobs=4)]: Done  17 tasks      | elapsed:   2.9min
[Parallel(n_jobs=4)]: Done  24 tasks      | elapsed:   7.2min
[Parallel(n_jobs=4)]: Done  33 out of  36 | elapsed: 16.0min remaining:  1.4min
[Parallel(n_jobs=4)]: Done  36 out of  36 | elapsed: 18.8min finished
```

```
Out[56]: GridSearchCV(cv=4, error_score='raise-deprecating',
                      estimator=LogisticRegression(C=1.0, class_weight='balanced', dual=False,
                                                    fit_intercept=True, intercept_scaling=1, max_iter=100,
                                                    multi_class='warn', n_jobs=None, penalty='l2', random_state=None,
                                                    solver='warn', tol=0.0001, verbose=0, warm_start=False),
                      fit_params=None, iid='warn', n_jobs=4,
                      param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]}},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring='roc_auc', verbose=10)
```

```
In [57]: k=a
         auc_cv=clf.cv_results_['mean_test_score']
         auc_train=clf.cv_results_['mean_train_score']
         plt.plot(np.log(k),auc_cv)
         plt.plot(np.log(k),auc_train)
         plt.title('C vs AUC score')
         plt.xlabel('C of Logistic regression')
         plt.ylabel('auc score')
         plt.legend({"test":"","train":""})
```

```
Out[57]: <matplotlib.legend.Legend at 0x2681315d080>
```



```
In [61]: print(clf.cv_results_['mean_test_score'])
         print(np.log(k))
```

```
[0.70631602 0.72820106 0.71590037 0.69237698 0.6781808  0.67466286
 0.67347534 0.67518525 0.66187874]
[-9.21034037 -6.90775528 -4.60517019 -2.30258509  0.          2.30258509
 4.60517019  6.90775528  9.21034037]
```

```
In [62]: print(np.exp(-6.90775528))
```

```
0.0009999999989821371
```

10**-3 is the optimal value

```
In [63]: model=LogisticRegression(class_weight='balanced',C=10**-3,n_jobs=4)
         model.fit(BOW,project_data_Y_train)
```

```
Out[63]: LogisticRegression(C=0.001, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=4, penalty='l2', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

```

In [64]: #https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classi
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from tqdm import tqdm

probs_test = model.predict_proba(BOW_test)
# keep probabilities for the positive outcome only
probs_test = probs_test[:, 1]
auc_test = roc_auc_score(project_data_Y_test, probs_test)
print('AUC: %.3f' % auc_test)
fpr, tpr, thresholds = roc_curve(project_data_Y_test, probs_test)


probs_train = model.predict_proba(BOW)
# keep probabilities for the positive outcome only
probs_train = probs_train[:, 1]
auc_train = roc_auc_score(project_data_Y_train, probs_train)
print('AUC: %.3f' % auc_train)
fpr1, tpr1, thresholds1 = roc_curve(project_data_Y_train, probs_train)

plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr1, tpr1, marker='.')
plt.plot(fpr, tpr, marker='.')

plt.legend({"standard": "", "train auc": "", "test auc": ""})
plt.title("AUC for tain and test using bag of words")
plt.xlabel("false positive rate")
plt.ylabel("true positive rate")
plt.show()

AUC: 0.736
AUC: 0.772

```




```
In [65]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(BOW_test)
tn, fp, fn, tp = confusion_matrix(project_data_Y_test,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate", (tp/(tp+fn)))
print("true negaitive rate", (tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("test confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

2243 1073 5868 12666

true positive rate 0.683392683716413

true negaitive rate 0.6764173703256936

```
[[2243, 5868], [1073, 12666]]
```




```
In [66]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(BOW)
tn, fp, fn, tp = confusion_matrix(project_data_Y_train,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate", (tp/(tp+fn)))
print("true negaitive rate", (tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("train confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

```

9564 3662 22889 51283
true positive rate 0.6914064606590088
true negative rate 0.7231211250567064
[[9564, 22889], [3662, 51283]]

```



2.0.2 2.4.1 Applying LR on TFIDF, SET 2

```

In [67]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack
# with the same hstack function we are concatenating a sparse matrix and a dense matrix
TFIDF = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_state_one_hot_train))
print(TFIDF.shape)
TFIDF_test = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_state_one_hot_test))
print(TFIDF_test.shape)

```

```

(87398, 7900)
(21850, 7900)

```

```

In [68]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
model=LogisticRegression(class_weight='balanced')

```

```

a=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
print(a)
parameters = {'C': a }
clf = GridSearchCV(model, parameters, cv=4,scoring='roc_auc',n_jobs=4,verbose=10)
clf.fit(TFIDF,project_data_Y_train)

```

```

[0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]
Fitting 4 folds for each of 9 candidates, totalling 36 fits

```

```

[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done   5 tasks      | elapsed:    9.5s
[Parallel(n_jobs=4)]: Done  10 tasks      | elapsed:   22.3s
[Parallel(n_jobs=4)]: Done  17 tasks      | elapsed:   1.1min
[Parallel(n_jobs=4)]: Done  24 tasks      | elapsed:   2.1min
[Parallel(n_jobs=4)]: Done  33 out of  36 | elapsed:   5.8min remaining:   31.8s
[Parallel(n_jobs=4)]: Done  36 out of  36 | elapsed:   6.6min finished

```

```

Out[68]: GridSearchCV(cv=4, error_score='raise-deprecating',
    estimator=LogisticRegression(C=1.0, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, max_iter=100,
    multi_class='warn', n_jobs=None, penalty='l2', random_state=None,
    solver='warn', tol=0.0001, verbose=0, warm_start=False),
    fit_params=None, iid='warn', n_jobs=4,
    param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='roc_auc', verbose=10)

```

```


In [69]: k=a
    auc_cv=clf.cv_results_['mean_test_score']
    auc_train=clf.cv_results_['mean_train_score']
    plt.plot(np.log(k),auc_cv)
    plt.plot(np.log(k),auc_train)
    plt.title('C vs AUC score')
    plt.xlabel('C of Logistic regression')
    plt.ylabel('auc score')
    plt.legend({"test":"","train":""})

```

```

Out[69]: <matplotlib.legend.Legend at 0x2681eccc978>

```

```
In [70]: print(clf.cv_results_['mean_test_score'])
         print(np.log(k))
```

```
[0.61549793 0.63605198 0.69201184 0.72625863 0.71027254 0.68017441
 0.6679054  0.66519146 0.66432133]
[-9.21034037 -6.90775528 -4.60517019 -2.30258509  0.          2.30258509
 4.60517019  6.90775528  9.21034037]
```

```
In [71]: print(np.exp(-2.30258509))
```

```
0.10000000029940456
```

10**-1 is the optimal value

```
In [72]: model=LogisticRegression(class_weight='balanced',C=10**-1,n_jobs=4)
         model.fit(TFIDF,project_data_Y_train)
```

```
Out[72]: LogisticRegression(C=0.1, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=4, penalty='l2', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

```

In [73]: #https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classi
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from tqdm import tqdm

probs_test = model.predict_proba(TFIDF_test)
# keep probabilities for the positive outcome only
probs_test = probs_test[:, 1]
auc_test = roc_auc_score(project_data_Y_test, probs_test)
print('AUC: %.3f' % auc_test)
fpr, tpr, thresholds = roc_curve(project_data_Y_test, probs_test)


probs_train = model.predict_proba(TFIDF)
# keep probabilities for the positive outcome only
probs_train = probs_train[:, 1]
auc_train = roc_auc_score(project_data_Y_train, probs_train)
print('AUC: %.3f' % auc_train)
fpr1, tpr1, thresholds1 = roc_curve(project_data_Y_train, probs_train)

plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr1, tpr1, marker='.')
plt.plot(fpr, tpr, marker='.')

plt.legend({"standard": "", "train auc": "", "test auc": ""})
plt.title("AUC for tain and test using bag of words")
plt.xlabel("false positive rate")
plt.ylabel("true positive rate")
plt.show()

AUC: 0.735
AUC: 0.789


```



```
In [74]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(TFIDF_test)
tn, fp, fn, tp = confusion_matrix(project_data_Y_test,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate",(tp/(tp+fn)))
print("true negaitive rate",(tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("test confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

```
2191 1125 5669 12865
true positive rate 0.6941297075644761
true negaitive rate 0.6607358262967431
```

```
[[2191, 5669], [1125, 12865]]
```




```
In [75]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(TFIDF)
tn, fp, fn, tp = confusion_matrix(project_data_Y_train,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate", (tp/(tp+fn)))
print("true negaitive rate", (tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("train confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

```

9683 3543 21773 52399
true positive rate 0.7064525697028529
true negative rate 0.7321185543626191
[[9683, 21773], [3543, 52399]]

```



2.0.3 2.4.1 Applying LR on average word to vector, SET 3

```

In [76]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack
# with the same hstack function we are concatenating a sparse matrix and a dense matrix
AVG_W2V = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_state_train))
print(AVG_W2V.shape)
AVG_W2V_test = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_state_test))
print(AVG_W2V_test.shape)

```

```

(87398, 697)
(21850, 697)

```

```

In [77]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
model=LogisticRegression(class_weight='balanced')

```

```

a=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
print(a)
parameters = {'C': a }
clf = GridSearchCV(model, parameters, cv=4,scoring='roc_auc',n_jobs=4,verbose=10)
clf.fit(AVG_W2V,project_data_Y_train)

```

```

[0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]
Fitting 4 folds for each of 9 candidates, totalling 36 fits

```

```

[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done   5 tasks      | elapsed:   46.3s
[Parallel(n_jobs=4)]: Done  10 tasks      | elapsed:   1.6min
[Parallel(n_jobs=4)]: Done  17 tasks      | elapsed:   5.0min
[Parallel(n_jobs=4)]: Done  24 tasks      | elapsed:   9.1min
[Parallel(n_jobs=4)]: Done  33 out of  36 | elapsed: 17.1min remaining:  1.6min
[Parallel(n_jobs=4)]: Done  36 out of  36 | elapsed: 18.0min finished

```

```

Out[77]: GridSearchCV(cv=4, error_score='raise-deprecating',
    estimator=LogisticRegression(C=1.0, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, max_iter=100,
    multi_class='warn', n_jobs=None, penalty='l2', random_state=None,
    solver='warn', tol=0.0001, verbose=0, warm_start=False),
    fit_params=None, iid='warn', n_jobs=4,
    param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='roc_auc', verbose=10)

```

```

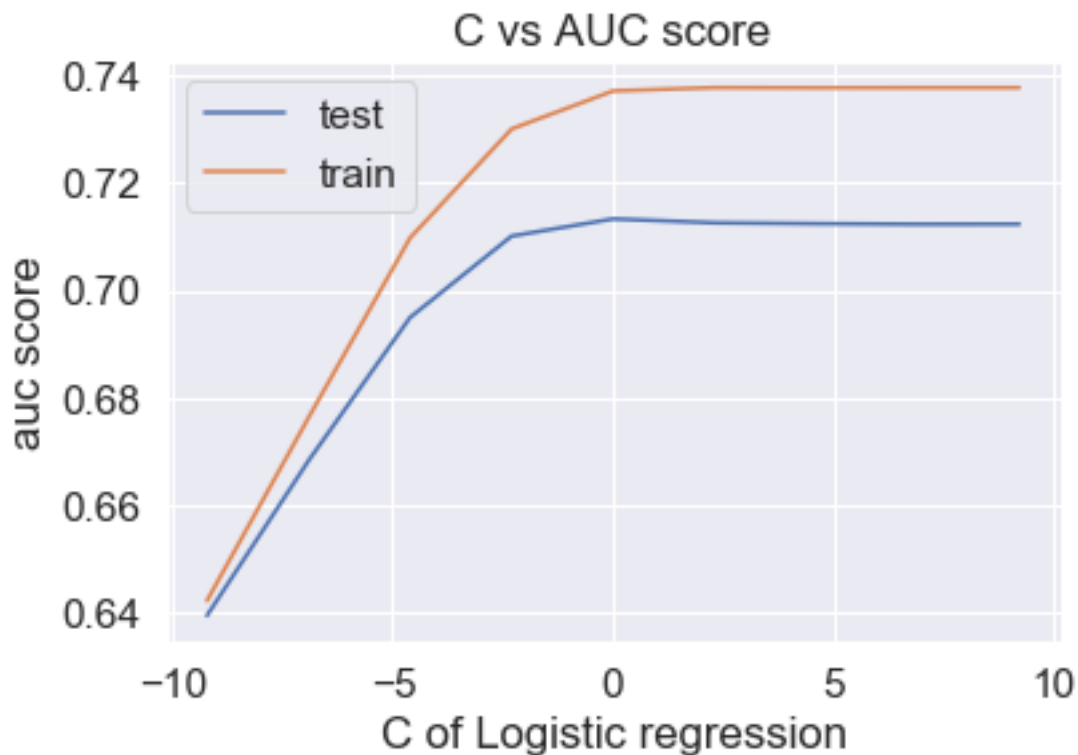
In [78]: k=a
    auc_cv=clf.cv_results_['mean_test_score']
    auc_train=clf.cv_results_['mean_train_score']
    plt.plot(np.log(k),auc_cv)
    plt.plot(np.log(k),auc_train)
    plt.title('C vs AUC score')
    plt.xlabel('C of Logistic regression')
    plt.ylabel('auc score')
    plt.legend({"test":"","train":""})

```

```

Out[78]: <matplotlib.legend.Legend at 0x2680e01ed30>

```



```
In [79]: print(clf.cv_results_['mean_test_score'])
         print(np.log(k))
```

```
[0.63930371 0.66821107 0.6949272  0.71011909 0.71332398 0.71259163
 0.71243786 0.71233063 0.71234299]
[-9.21034037 -6.90775528 -4.60517019 -2.30258509  0.          2.30258509
 4.60517019  6.90775528  9.21034037]
```

```
In [80]: print(np.exp(0))
```

```
1.0
```

10**0 is the optimal value

```
In [81]: model=LogisticRegression(class_weight='balanced',C=1,n_jobs=4)
         model.fit(AVG_W2V,project_data_Y_train)
```

```
Out[81]: LogisticRegression(C=1, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=4, penalty='l2', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

```

In [82]: #https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classi
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from tqdm import tqdm

probs_test = model.predict_proba(AVG_W2V_test)
# keep probabilities for the positive outcome only
probs_test = probs_test[:, 1]
auc_test = roc_auc_score(project_data_Y_test, probs_test)
print('AUC: %.3f' % auc_test)
fpr, tpr, thresholds = roc_curve(project_data_Y_test, probs_test)

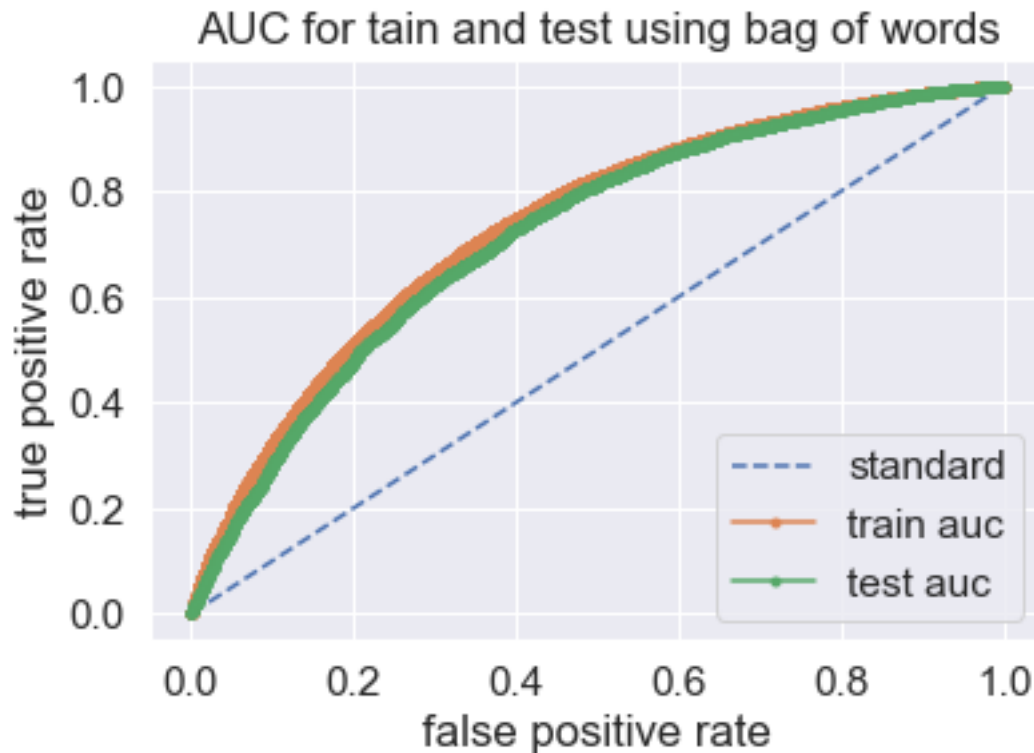
probs_train = model.predict_proba(AVG_W2V)
# keep probabilities for the positive outcome only
probs_train = probs_train[:, 1]
auc_train = roc_auc_score(project_data_Y_train, probs_train)
print('AUC: %.3f' % auc_train)
fpr1, tpr1, thresholds1 = roc_curve(project_data_Y_train, probs_train)

plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr1, tpr1, marker='.')
plt.plot(fpr, tpr, marker='.')

plt.legend({"standard": "", "train auc": "", "test auc": ""})
plt.title("AUC for tain and test using bag of words")
plt.xlabel("false positive rate")
plt.ylabel("true positive rate")
plt.show()

AUC: 0.718
AUC: 0.734


```

```
In [83]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(AVG_W2V_test)
tn, fp, fn, tp = confusion_matrix(project_data_Y_test,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate",(tp/(tp+fn)))
print("true negaitive rate",(tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("test confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

```
2204 1112 6308 12226
true positive rate 0.659652530484515
true negaitive rate 0.6646562123039808
```

[[2204, 6308], [1112, 12226]]




```
In [84]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(AVG_W2V)
tn, fp, fn, tp = confusion_matrix(project_data_Y_train,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate", (tp/(tp+fn)))
print("true negaitive rate", (tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("train confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

```

9028 4198 24862 49310
true positive rate 0.6648061263010301
true negative rate 0.6825948888552851
[[9028, 24862], [4198, 49310]]

```



2.0.4 2.4.1 Applying LR on tfidf word to vector, SET 4

```

In [85]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack
# with the same hstack function we are concatenating a sparse matrix and a dense matrix
TFIDF_W2V = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_status_train))
print(TFIDF_W2V.shape)
TFIDF_W2V_test = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_status_test))
print(TFIDF_W2V_test.shape)

```

```

(87398, 697)
(21850, 697)

```

```

In [86]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
model=LogisticRegression(class_weight='balanced')

```

```

a=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
print(a)
parameters = {'C': a }
clf = GridSearchCV(model, parameters, cv=4,scoring='roc_auc',n_jobs=4,verbose=10)
clf.fit(TFIDF_W2V,project_data_Y_train)

```

```

[0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]
Fitting 4 folds for each of 9 candidates, totalling 36 fits

```

```

[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done   5 tasks      | elapsed:   44.8s
[Parallel(n_jobs=4)]: Done  10 tasks      | elapsed:   1.5min
[Parallel(n_jobs=4)]: Done  17 tasks      | elapsed:   4.2min
[Parallel(n_jobs=4)]: Done  24 tasks      | elapsed:   6.1min
[Parallel(n_jobs=4)]: Done  33 out of  36 | elapsed: 11.4min remaining:  1.0min
[Parallel(n_jobs=4)]: Done  36 out of  36 | elapsed: 11.9min finished

```

```

Out[86]: GridSearchCV(cv=4, error_score='raise-deprecating',
    estimator=LogisticRegression(C=1.0, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, max_iter=100,
    multi_class='warn', n_jobs=None, penalty='l2', random_state=None,
    solver='warn', tol=0.0001, verbose=0, warm_start=False),
    fit_params=None, iid='warn', n_jobs=4,
    param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='roc_auc', verbose=10)

```

```

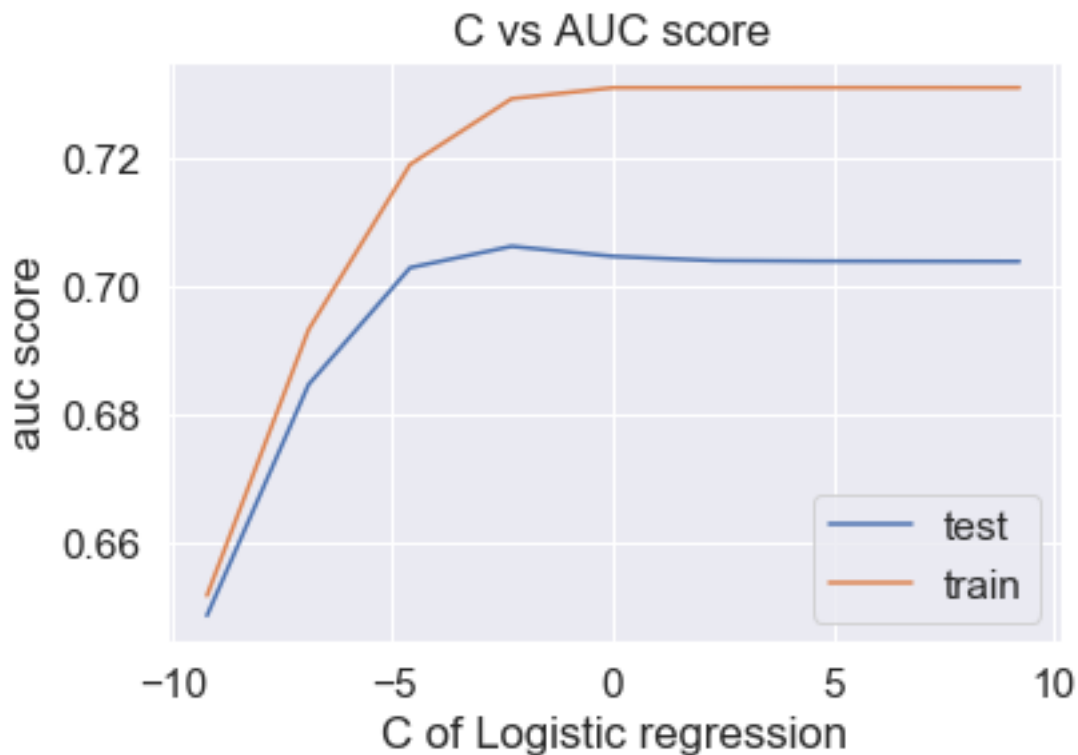
In [87]: k=a
    auc_cv=clf.cv_results_['mean_test_score']
    auc_train=clf.cv_results_['mean_train_score']
    plt.plot(np.log(k),auc_cv)
    plt.plot(np.log(k),auc_train)
    plt.title('C vs AUC score')
    plt.xlabel('C of Logistic regression')
    plt.ylabel('auc score')
    plt.legend({"test":"","train":""})

```

```

Out[87]: <matplotlib.legend.Legend at 0x2682f66e9e8>

```



```
In [88]: print(clf.cv_results_['mean_test_score'])
         print(np.log(k))
```

```
[0.64873826 0.68464745 0.70284485 0.70616459 0.7046037  0.70394507
 0.70385525 0.70383068 0.70380748]
[-9.21034037 -6.90775528 -4.60517019 -2.30258509  0.          2.30258509
 4.60517019  6.90775528  9.21034037]
```

```
In [89]: print(np.exp(-2.30258509))
```

```
0.10000000029940456
```

10**-1 is the optimal value

```
In [90]: model=LogisticRegression(class_weight='balanced',C=10**-1,n_jobs=4)
         model.fit(TFIDF_W2V,project_data_Y_train)
```

```
Out[90]: LogisticRegression(C=0.1, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=4, penalty='l2', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

```

In [91]: #https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classi
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from tqdm import tqdm

probs_test = model.predict_proba(TFIDF_W2V_test)
# keep probabilities for the positive outcome only
probs_test = probs_test[:, 1]
auc_test = roc_auc_score(project_data_Y_test, probs_test)
print('AUC: %.3f' % auc_test)
fpr, tpr, thresholds = roc_curve(project_data_Y_test, probs_test)


probs_train = model.predict_proba(TFIDF_W2V)
# keep probabilities for the positive outcome only
probs_train = probs_train[:, 1]
auc_train = roc_auc_score(project_data_Y_train, probs_train)
print('AUC: %.3f' % auc_train)
fpr1, tpr1, thresholds1 = roc_curve(project_data_Y_train, probs_train)

plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr1, tpr1, marker='.')
plt.plot(fpr, tpr, marker='.')

plt.legend({"standard": "", "train auc": "", "test auc": ""})
plt.title("AUC for tain and test using bag of words")
plt.xlabel("false positive rate")
plt.ylabel("true positive rate")
plt.show()

AUC: 0.711
AUC: 0.727


```



```
In [92]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(TFIDF_W2V_test)
tn, fp, fn, tp = confusion_matrix(project_data_Y_test,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate",(tp/(tp+fn)))
print("true negaitive rate",(tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("test confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

```
2214 1102 6678 11856
true positive rate 0.6396892198122369
true negaitive rate 0.6676718938480096
```

```
[[2214, 6678], [1102, 11856]]
```




```
In [93]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(TFIDF_W2V)
tn, fp, fn, tp = confusion_matrix(project_data_Y_train,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate", (tp/(tp+fn)))
print("true negaitive rate", (tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("train confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```



```

9166 4060 26430 47742
true positive rate 0.6436660734508979
true negative rate 0.6930288825041585
[[9166, 26430], [4060, 47742]]

```



2.5 Logistic Regression with added Features Set 5

```

In [94]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack
# with the same hstack function we are concatenating a sparse matrix and a dense matrix
WITHOUT_WORDS = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_type_train))
print(WITHOUT_WORDS.shape)
WITHOUT_WORDS_test = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_type_test))
print(WITHOUT_WORDS_test.shape)

```

```

(87398, 101)
(21850, 101)

```

```

In [95]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
model=LogisticRegression(class_weight='balanced')
a=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]

```

```

print(a)
parameters = {'C': a }
clf = GridSearchCV(model, parameters, cv=4,scoring='roc_auc',n_jobs=4,verbose=10)
clf.fit(WITHOUT_WORDS,project_data_Y_train)

```

```

[0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]
Fitting 4 folds for each of 9 candidates, totalling 36 fits

```

```

[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done   5 tasks      | elapsed:    1.5s
[Parallel(n_jobs=4)]: Done  10 tasks      | elapsed:    2.6s
[Parallel(n_jobs=4)]: Done  17 tasks      | elapsed:    5.6s
[Parallel(n_jobs=4)]: Done  24 tasks      | elapsed:    7.7s
[Parallel(n_jobs=4)]: Done  33 out of  36 | elapsed:   11.8s remaining:    1.0s
[Parallel(n_jobs=4)]: Done  36 out of  36 | elapsed:   12.4s finished

```

```

Out[95]: GridSearchCV(cv=4, error_score='raise-deprecating',
      estimator=LogisticRegression(C=1.0, class_weight='balanced', dual=False,
      fit_intercept=True, intercept_scaling=1, max_iter=100,
      multi_class='warn', n_jobs=None, penalty='l2', random_state=None,
      solver='warn', tol=0.0001, verbose=0, warm_start=False),
      fit_params=None, iid='warn', n_jobs=4,
      param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]}},
      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
      scoring='roc_auc', verbose=10)

```

```


In [96]: k=a
auc_cv=clf.cv_results_['mean_test_score']
auc_train=clf.cv_results_['mean_train_score']
plt.plot(np.log(k),auc_cv)
plt.plot(np.log(k),auc_train)
plt.title('C vs AUC score')
plt.xlabel('C of Logistic regression')
plt.ylabel('auc score')
plt.legend({"test":"","train":""})

```

```

Out[96]: <matplotlib.legend.Legend at 0x268331089e8>

```



```
In [97]: print(clf.cv_results_['mean_test_score'])
         print(np.log(k))
```

```
[0.61191807 0.61467438 0.61727165 0.61580968 0.61504122 0.61491431
 0.61496344 0.61489287 0.61490382]
[-9.21034037 -6.90775528 -4.60517019 -2.30258509  0.          2.30258509
 4.60517019  6.90775528  9.21034037]
```

```
In [98]: print(np.exp(-4.60517019))
```

```
0.0099999999959880915
```

10^{-1} is the optimal value

```
In [99]: model=LogisticRegression(class_weight='balanced',C=10**-2,n_jobs=4)
         model.fit(WITHOUT_WORDS,project_data_Y_train)
```

```
Out[99]: LogisticRegression(C=0.01, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=4, penalty='l2', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

```

In [100]: #https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-class
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from tqdm import tqdm

probs_test = model.predict_proba(WITHOUT_WORDS_test)
# keep probabilities for the positive outcome only
probs_test = probs_test[:, 1]
auc_test = roc_auc_score(project_data_Y_test, probs_test)
print('AUC: %.3f' % auc_test)
fpr, tpr, thresholds = roc_curve(project_data_Y_test, probs_test)


probs_train = model.predict_proba(WITHOUT_WORDS)
# keep probabilities for the positive outcome only
probs_train = probs_train[:, 1]
auc_train = roc_auc_score(project_data_Y_train, probs_train)
print('AUC: %.3f' % auc_train)
fpr1, tpr1, thresholds1 = roc_curve(project_data_Y_train, probs_train)

plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr1, tpr1, marker='.')
plt.plot(fpr, tpr, marker='.')

plt.legend({"standard": "", "train auc": "", "test auc": ""})
plt.title("AUC for tain and test using bag of words")
plt.xlabel("false positive rate")
plt.ylabel("true positive rate")
plt.show()

AUC: 0.624
AUC: 0.622


```



```
In [101]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(WITHOUT_WORDS_test)
tn, fp, fn, tp = confusion_matrix(project_data_Y_test,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate",(tp/(tp+fn)))
print("true negaitive rate",(tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("test confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

```
1970 1346 7718 10816
true positive rate 0.5835761303550232
true negaitive rate 0.5940892641737032
```

```
[[1970, 7718], [1346, 10816]]
```




```
In [102]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#compute confudion matrix values and plot
from sklearn.metrics import confusion_matrix
predicted_bow_test=model.predict(WITHOUT_WORDS)
tn, fp, fn, tp = confusion_matrix(project_data_Y_train,predicted_bow_test).ravel()
print(tn, fp, fn, tp)
print("true positive rate",(tp/(tp+fn)))
print("true negaitive rate",(tn/(tn+fp)))
matrix=[[tn,fn],[fp,tp]]
print(matrix)
df_cm = pd.DataFrame(matrix, range(2),
                      range(2))
#plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')# font size
plt.title("train confusion matrix using bag of words")
plt.xlabel("predicted")
plt.ylabel("actual")
plt.show()
```

```

7773 5453 30675 43497
true positive rate 0.5864342339427278
true negative rate 0.5877060335702404
[[7773, 30675], [5453, 43497]]

```



Without words the model is not performing well. The auc is 0.62 and the tpr and fpr are also not so good. Hence I reject this model

3. Conclusion

```

In [103]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Vectorizer", "Model", "alpha", "AUC"]
x.add_row(["BAG of words", "Logistic regression", 10**-3, 0.736])
x.add_row(["TFIDF", "Logistic regression", 10**-1, 0.735])
x.add_row(["Average W2V", "Logistic regression", 1, 0.718])
x.add_row(["TFIDF W2V", "Logistic regression", 10**-1, 0.711])
x.add_row(["Logistic Regression with added Features", "Logistic regression", 10**-1, 0.711])
x.border=True
print(x)

```

Vectorizer	Model	alpha	AUC
BAG of words	Logistic regression	10 ⁻³	0.736
TFIDF	Logistic regression	10 ⁻¹	0.735
Average W2V	Logistic regression	1	0.718
TFIDF W2V	Logistic regression	10 ⁻¹	0.711
Logistic Regression with added Features	Logistic regression	10 ⁻¹	0.711

+-----+-----+-----+-----+				
	BAG of words	Logistic regression	0.001	0.736
	TFIDF	Logistic regression	0.1	0.735
	Average W2V	Logistic regression	1	0.718
	TFIDF W2V	Logistic regression	0.1	0.711
	Logistic Regression with added Features	Logistic regression	0.1	0.624
+-----+-----+-----+-----+				